

# Exploring HBase 2

**Objective:** Use the HBase Shell to populate and explore the `hbase:meta` table.

**HBase tables:** `[some name]` (use your own table name)

In this exercise you will use the HBase Shell to explore the `hbase:meta` table.

Hbase is not a relational database. The operations available for data stored in hbase are the following:

- Get: retrieves a row or a subset of a row
- Put: Add or update a row or a subset of a row
- Scan: retrieve a range of sequential rows
- Delete: remove a row or a subset of a row

---

## Data Definition Language

`ddl` stands for Data Definition Language. So help on creating tables defining column family attributes will be displayed by typing:

```
hbase> help 'ddl'
```

```
hbase> help 'get'
```

```
hbase> help 'put'
```

```
hbase> help 'scan'  
hbase> help 'delete'
```

So let's do a couple of these.

```
  
<h4>1. Create the Table</h4>
```

Enter HBase Shell:

```
hbase shell  
hbase(main):001:0>
```

Create a table in HBase (again, name it your choice) and specify a column name:

```
hbase(main):007:0> create '[some name]', 'address'
```

Now check the tables:

```
hbase(main):008:0* list
```

Verify the new table:

```

hbase(main):012:0> describe '[some name]'
DESCRIPTION          ENABLED
'[some name]', {NAME => 'address', DATA_BLOCK_ENCODING =>
'NONE', BLOOMFILTER => 'ROW', REPLICATI true ON_SCOPE => '
0', VERSIONS => '1', COMPRESSION => 'NONE', MIN_VERSIONS =
> '0', TTL => '2147483647', KEEP_DELETED_CELLS => 'false',
BLOCKSIZE => '65536', IN_MEMORY => 'false', BLOCKCACHE =>
'true'}
1 row(s) in 0.0750 seconds

```

```


<h4>2. PUT Something in the Table</h4>

```

Review the output and notice the column ('address'), parameter properties such as COMPRESSION, DATA\_BLOCK\_ENCODING, whether is kept in memory...

Insert data using PUT command :

```

hbase(main):004:0> put '[some name]','row1','address','val
ue1'
0 row(s) in 0.4250 seconds
hbase(main):005:0> put '[some name]','row2','address','val
ue1'

```

```
0 row(s) in 0.0120 seconds
```

```
hbase(main):006:0> put '[some name]','row1','address','value2'
```

```
0 row(s) in 0.0110 seconds
```

```
hbase(main):007:0> put '[some name]','row1','address','value3'
```

```
0 row(s) in 0.0070 seconds
```

View the table data using SCAN:

```
hbase(main):008:0> scan '[some name]'
```

ROW	COLUMN+CELL
-----	-------------

row1	column=address:, timestamp=1408499993409, value=value3
------	--

row2	column=address:, timestamp=1408499977588, value=value1
------	--

```
2 row(s) in 0.1020 seconds
```

This time view the data using SCAN with versions:

```
hbase(main):011:0> scan '[some name]', { VERSIONS => 3}
```

ROW	COLUMN+CELL
-----	-------------

row1	column=address:, timestamp=1408499993409, value=value3
------	--

row2	column=address:, timestamp=1408499977588, value=value1
------	--

```
2 row(s) in 0.0200 seconds
```

Use the **count** command to determine the number of rows in a table:

```
hbase(main):002:0> count '[some name]'
```

2 row(s) in 0.4020 seconds
----------------------------

```
=> 2
```



<h4>3. Retrieve data using GET</h4>

```
hbase(main):012:0> get '[some name]','row1'
```

COLUMN	CELL
address:	timestamp=1408499993409, value=value3

```
1 row(s) in 0.0530 seconds
```

Use GET command to check on row2:

```
hbase(main):003:0> get '[some name]','row2'
```

COLUMN	CELL
address:	timestamp=1408499977588, value=value1

```
1 row(s) in 0.0110 seconds
```



<h4>4. Now use PUT command on row2</h4>

```
hbase(main):004:0> put '[some name]','row2','address','value2'  
0 row(s) in 0.3340 seconds
```

Retrieve the value:

```
hbase(main):005:0> get '[some name]','row2'
```

Test the row counts:

```
hbase(main):006:0> count '[some name]'  
2 row(s) in 0.2290 seconds  
=> 2
```

<img alt="ToDo Logo" data-bbox="61 708 125 765"/> [images.githubusercontent.com/558905/40613898-7a6c70d6-624e-11e8-9178-7bde851ac7bd.png](https://user-images.githubusercontent.com/558905/40613898-7a6c70d6-624e-11e8-9178-7bde851ac7bd.png) align="left" width="50" height="50"  
title="ToDo Logo">  
<h4>5. Run the delete command</h4>

```
hbase(main):020:0> delete '[some name]','row1','address'
```

```
0 row(s) in 0.1990 seconds
```

We can verify the data using SCAN command:

```
hbase(main):021:0> scan '[some name]'
```

ROW	COLUMN+CELL
row2	column=address:, timestamp=1409577011732, value=value2

```
1 row(s) in 0.0750 seconds
```

```

<h4>6. Disable</h4>
```

Before we disable the table, let's run a count on the table:

```
hbase(main):022:0> count '[some name]'
```

```
1 row(s) in 0.0320 seconds
```

```
=> 1
```

Execute the disable now:

```
hbase(main):007:0> disable '[some name]'
```

```
0 row(s) in 1.7060 seconds
```

Do the count command again:

```
hbase(main):008:0> count '[some name]'
```

Now enable the table again:

```
hbase(main):025:0> enable '[some name]'
```

```
0 row(s) in 0.8730 seconds
```

Use count to check the table again:

```
hbase(main):026:0> count '[some name]'
```

Check the table again:

```
hbase(main):010:0> list
```

TABLE
...
[some name]
...



<h4>7. Drop</h4>



Now use the DROP command to get rid of the table:

```
hbase(main):011:0> drop '[some name]'  
0 row(s) in 0.8290 seconds
```

To verify use the list command again:

```
hbase(main):012:0> list
```

## Results

You're getting the hang of it! Good!

<https://virtuant.github.io/hadoop-overview-spark-hwx/>>Go Back

<br>

<br>