

Lab: HBase and Hive Integration

Objective: Understand how HBase and Hive integrate. You will complete data storage in HBase from Hive table data.

Data directory: `~/data/hbase/data`

What opportunities exist for deeper integration? Currently, customers are putting together solutions leveraging HBase, Phoenix, Hive etc. to build bespoke a closed-loop system for operational data and SQL analytics. We feel there is an opportunity to provide out-of-the-box integration with ease of use and additional capabilities such as transactions, cross datacenter failover etc.

So let's take a look at Hive -> HBase integration.

What is ACID?

ACID stands for:

- Atomicity: a transaction should complete successfully or else it should fail completely i.e. it should not be left partially
- Consistency: any transaction will bring the database from one valid state to another state
- Isolation: every transaction should be independent of each other i.e. one transaction should not affect another
- Durability: if a transaction is completed, it should be preserved

in the database even if the machine state is lost or a system failure might occur

These ACID properties are essential for a transaction and every transaction should ensure that these properties are met.

Transactions in Hive

Transactions in Hive are introduced in Hive 0.13, but they only partially fulfill the ACID properties like atomicity, consistency, durability, at the partition level. Here, Isolation can be provided by turning on one of the locking mechanisms available with zookeeper or in memory.

But in Hive 0.14, new API's have been added to completely fulfill the ACID properties while performing any transaction.

Transactions are provided at the row-level in Hive 0.14. The different row-level transactions available in Hive are as follows:

- Insert
- Delete
- Update

There are numerous limitations with the present transactions available in Hive.

ORC is the file format supported by Hive transaction. It is now essential to have ORC file format for performing transactions in Hive. The table

needs to be bucketed in order to support transactions.

Missing Files

Make sure that the node you've picked is available to run **beeline**:

```
beeline -n hive -u jdbc:hive2://localhost:10000
```

You may see something like this:

```
Connecting to jdbc:hive2://localhost:10000
18/09/25 03:56:01 [main]: WARN jdbc.HiveConnection: Failed
to connect to localhost:10000
Could not open connection to the HS2 server. Please check
the server URI and if the URI is correct, then ask the adm
inistrator to check the server status.
Error: Could not open client transport with JDBC Uri: jdbc
:hive2://localhost:10000: java.net.ConnectException: Conne
ction refused (Connection refused) (state=08S01,code=0)
Beeline version 3.1.0.3.0.1.0-187 by Apache Hive
[centos@ip-172-30-12-85 ~]$
```

If so, try another node as a Hive client. Once you find a running node, you may need to copy the files from **data** on the Ambari (or other) node to your local drive. You can do an **scp** (a SSH copy of files) from the Ambari node like this:

```
[centos@ip-172-30-13-166 ~]$ sudo su -  
[root@ip-172-30-13-166 ~]# scp -r /home/centos/data root@[  
remote Hive node address]:/home/centos
```

Note: scp works the other way too

Go to the Hive node, and you should now see your files:

```
[centos@ip-172-30-11-227 ~]$ ll  
total 0  
drwxrwxr-x. 5 root root 51 Sep 25 03:40 data
```

It may require you to **chown** the **data** directory as well:

```
[centos@ip-172-30-11-227 ~]$ sudo chown -R centos:centos d  
ata/
```

Look at the Data

Do a **more** to look at the data in **iot_data.csv**:

```
[centos@ip-10-0-0-237 data]$ more iot_data.csv  
_id,deviceParameter,deviceValue,deviceId,dateTime  
ObjectId(5a81b5395882b86112555f70),Temperature,27,SBS0  
5,2018-02-12 15:39:37.050 UTC  
ObjectId(5a81b5395882b86112555f71),Humidity,59,SBS05,2
```

```
018-02-12 15:39:37.801 UTC
```

```
    ObjectId(5a81b53a5882b86112555f72) , Sound , 130 , SBS04 , 201
```

```
8-02-12 15:39:38.629 UTC
```

```
    ObjectId(5a81b53b5882b86112555f73) , Humidity , 75 , SBS05 , 2
```

```
018-02-12 15:39:39.272 UTC
```

```
    ObjectId(5a81b53b5882b86112555f74) , Temperature , 33 , SBS0
```

```
2 , 2018-02-12 15:39:39.613 UTC
```

```
    ObjectId(5a81b53c5882b86112555f75) , Sound , 102 , SBS03 , 201
```

```
8-02-12 15:39:40.363 UTC
```

```
    ObjectId(5a81b53c5882b86112555f76) , Temperature , 18 , SBS0
```

```
2 , 2018-02-12 15:39:40.663 UTC
```

```
    ObjectId(5a81b53c5882b86112555f77) , Flow , 64 , SBS05 , 2018-
```

```
02-12 15:39:40.678 UTC
```

```
    ObjectId(5a81b53d5882b86112555f78) , Temperature , 28 , SBS0
```

```
4 , 2018-02-12 15:39:41.141 UTC
```

For bypassing any security issues, put the `iot_data.csv` table into the `tmp` directory in the Linux system:

```
cp iot_data.csv /tmp/.
```

Let's go ahead and create the anonymous user in HDFS:

```
sudo su hdfs
```

```
hdfs@host:~$ hdfs dfs -mkdir /user/anonymous
```

```
hdfs@host:~$ hdfs dfs -chown anonymous /user/anonymous
```

Create a Hive table and Load Data

We will begin by creating a Hive table with HBase characteristics. Notice the `hbase.table.name` below.

Now start beeline and create the table `iot_data`:

```
CREATE EXTERNAL TABLE iot_data
    (rowkey string, parameter string, value int, device_id string, datetime string)
    ROW FORMAT SERDE 'org.apache.hadoop.hive.hbase.HBaseSerDe'
    STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
    WITH SERDEPROPERTIES ("hbase.columns.mapping"=":key,para:parameter,para:value,para:device_id,para:datetime")
    TBLPROPERTIES ("hbase.table.name"="iot_data");
```

Easier to copy:

```
CREATE EXTERNAL TABLE iot_data (rowkey string, parameter string, value int, device_id string, datetime string) ROW FORMAT SERDE 'org.apache.hadoop.hive.hbase.HBaseSerDe' STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler' WITH SERDEPROPERTIES ("hbase.columns.mapping"=":key,para:parameter,para:value,para:device_id,para:datetime") TBLPROPERTIES ("hbase.table.name"="iot_data");
```

```
ERTIES ("hbase.table.name"="iot_data");
```

Validate the table in Hive:

```
describe iot_data;
INFO  : Compiling command(queryId=hive_20180815000016_8eed1209-306c-4e5b-8e19-45fb250fae0a): describe iot_data
INFO  : Semantic Analysis Completed (retrial = false)
INFO  : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:col_name, type:string, comment:from deserializer), FieldSchema(name:data_type, type:string, comment:from deserializer), FieldSchema(name:comment, type:string, comment:from deserializer)], properties:null)
INFO  : Completed compiling command(queryId=hive_20180815000016_8eed1209-306c-4e5b-8e19-45fb250fae0a); Time taken: 0.03 seconds
INFO  : Executing command(queryId=hive_20180815000016_8eed1209-306c-4e5b-8e19-45fb250fae0a): describe iot_data
INFO  : Starting task [Stage-0:DDL] in serial mode
INFO  : Completed executing command(queryId=hive_20180815000016_8eed1209-306c-4e5b-8e19-45fb250fae0a); Time taken: 0.026 seconds
INFO  : OK
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| id       | string   |         |
```

parameter	string		
value	int		
device_id	string		
datetime	string		

+-----+-----+-----+

5 rows selected (0.291 seconds)

Now we will create a temporary lookup table in Hive:

```
CREATE TABLE iot_in (
    id string, parameter string, value string, device_
id string, datetime string
) ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.
OpenCSVSerde'
WITH SERDEPROPERTIES ( "separatorChar" = ",", "quo
teChar" = "\"")
STORED AS TEXTFILE location '/tmp/iot_data.csv';
```

Easier to copy:

```
CREATE TABLE iot_in (id string, parameter string, value st
ring, device_id string, datetime string) ROW FORMAT SERDE
'org.apache.hadoop.hive.serde2.OpenCSVSerde' WITH SERDEPRO
PERTIES ( "separatorChar" = ",", "quoteChar" = "\"") STORE
D AS TEXTFILE location '/tmp/iot_data.csv';
```


And describe it:

```
describe iot_in;
```

and should be empty:

```
INFO  : OK

+-----+-----+-----+-----+
| iot_in.id | iot_in.parameter | iot_in.value | iot_in.device_id | iot_in.datetime |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+

No rows selected (0.121 seconds)
```

Now load data into the table:

```
LOAD DATA LOCAL INPATH '/tmp/iot_data.csv' OVERWRITE I
NTO TABLE iot_in;
```

Now print to make sure its loaded:

```
select * from iot_in limit 20;

INFO  : Returning Hive schema: Schema(fieldSchemas:[FieldS
```

```
chema(name:iot_in.id, type:int, comment:null), FieldSchema
(name:iot_in.parameter, type:string, comment:null), FieldS
chema(name:iot_in.value, type:int, comment:null), FieldSch
ema(name:iot_in.device_id, type:string, comment:null), Fie
ldSchema(name:iot_in.datetime, type:string, comment:null)]
, properties:null)
```

INFO : Completed compiling command(queryId=hive_20180815033623_dcb145d1-b7e7-4af6-b4bd-4c78b4f6ee15); Time taken: 0.209 seconds

INFO : Executing command(queryId=hive_20180815033623_dcb145d1-b7e7-4af6-b4bd-4c78b4f6ee15): `select * from iot_in limit 20`

INFO : Completed executing command(queryId=hive_20180815033623_dcb145d1-b7e7-4af6-b4bd-4c78b4f6ee15); Time taken: 0.001 seconds

INFO : OK

+-----+-----+-----+-----			
-----+-----+-----+-----			
iot_in.id	iot_in.parameter	iot_in.value	iot_in.
device_id	iot_in.datetime		
+-----+-----+-----+-----			
-----+-----+-----+-----			
ObjectId(4	deviceParameter	NULL	deviceI
d	dateTime		
ObjectId(3	Temperature	27	SBS05
	2018-02-12 15:39:37.050 UTC		
ObjectId(4	Humidity	59	SBS05
	2018-02-12 15:39:37.801 UTC		

	ObjectId(7		Sound		130		SBS04
			2018-02-12 15:39:38.629 UTC				
	ObjectId(6		Humidity		75		SBS05
			2018-02-12 15:39:39.272 UTC				
	ObjectId(9		Temperature		33		SBS02
			2018-02-12 15:39:39.613 UTC				
	ObjectId(8		Sound		102		SBS03
			2018-02-12 15:39:40.363 UTC				
	ObjectId(2		Temperature		18		SBS02
			2018-02-12 15:39:40.663 UTC				
	ObjectId(5		Flow		64		SBS05
			2018-02-12 15:39:40.678 UTC				
	ObjectId(6		Temperature		28		SBS04
			2018-02-12 15:39:41.141 UTC				
	ObjectId(1		Humidity		69		SBS03
			2018-02-12 15:39:41.804 UTC				
	ObjectId(8		Temperature		19		SBS04
			2018-02-12 15:39:42.350 UTC				
	ObjectId(6		Temperature		28		SBS05
			2018-02-12 15:39:42.593 UTC				
	ObjectId(6		Temperature		31		SBS04
			2018-02-12 15:39:43.070 UTC				
	ObjectId(2		Sound		133		SBS05
			2018-02-12 15:39:43.961 UTC				
	ObjectId(6		Flow		99		SBS02
			2018-02-12 15:39:44.031 UTC				
	ObjectId(4		Humidity		65		SBS04
			2018-02-12 15:39:44.667 UTC				

	ObjectId(3		Humidity		90		SBS05
			2018-02-12 15:39:45.260 UTC				
	ObjectId(4		Flow		89		SBS05
			2018-02-12 15:39:45.460 UTC				
	ObjectId(4		Sound		140		SBS03
			2018-02-12 15:39:46.389 UTC				

```
+-----+-----+-----+-----+
-----+-----+-----+-----+
```

20 rows selected (0.324 seconds)

Now, do a count:

```
0: jdbc:hive2://localhost:10000> select count(*) from iot_
in;
INFO  : Compiling command(queryId=hive_20180925055819_1077
c68f-8aa0-4f74-a7dd-e817f93bc3d7): select count(*) from io
t_in
INFO  : Semantic Analysis Completed (retrial = false)
INFO  : Returning Hive schema: Schema(fieldSchemas:[FieldS
chema(name:_c0, type:bigint, comment:null)], properties:nu
ll)
INFO  : Completed compiling command(queryId=hive_201809250
55819_1077c68f-8aa0-4f74-a7dd-e817f93bc3d7); Time taken: 0
.371 seconds
INFO  : Executing command(queryId=hive_20180925055819_1077
c68f-8aa0-4f74-a7dd-e817f93bc3d7): select count(*) from io
t_in
```

INFO : Query ID = hive_20180925055819_1077c68f-8aa0-4f74-a7dd-e817f93bc3d7

INFO : Total jobs = 1

INFO : Launching Job 1 out of 1

INFO : Starting task [Stage-1:MAPRED] in serial mode

INFO : Subscribed to counters: [] for queryId: hive_20180925055819_1077c68f-8aa0-4f74-a7dd-e817f93bc3d7

INFO : Session is already open

INFO : Dag name: select count(*) from iot_in (Stage-1)

INFO : Tez session was closed. Reopening...

INFO : Session re-established.

INFO : Session re-established.

INFO : Status: Running (Executing on YARN cluster with App id application_1537813688041_0010)

VERTICES		MODE	STATUS	TOTAL	COMPLETED
RUNNING	PENDING	FAILED	KILLED		

Map 1	container	SUCCEEDED	2	2
	0	0	0	0	
Reducer 2	container	SUCCEEDED	1	1
	0	0	0	0	

VERTICES: 02/02 [=====>>] 100% ELAP

```
SED TIME: 9.44 s
```

```
...
```

There is a ton of DAG and Map/Reduce action displayed here.

Question: how many records did the table have? 426,879?

Exit Hive by executing a `!q`:

```
!q
```

```
Closing: 0: jdbc:hive2://master1.hdp.com:2181/default;  
serviceDiscoveryMode=zooKeeper;zooKeeperNamespace=hiveserv  
er2
```

Validate the Table in HBase

Enter HBase shell:

```
hbase shell
```

Use LIST command to check tables:

```
hbase(main):005:0> list
```

```
TABLE
```

```
...
```

```
iot_data
```

```
iot_in
```

```
...
```

Validate the table in HBase:

```
hbase(main):003:0> describe 'iot_in'
```

```
Table iot_in is ENABLED
```

```
iot_in
```

```
COLUMN FAMILIES DESCRIPTION
```

```
{NAME => 'RAW', VERSIONS => '1', EVICT_BLOCKS_ON_CLOSE  
=> 'false', NEW_VERSION_BEHAVIOR => 'false', KEEP_DELETE  
D_CELLS => 'FALSE', CACHE_DATA_ON_WRITE => 'false', DA  
TA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIO  
NS => '0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'R  
OW', CACHE_INDEX_ON_WRITE => 'false', IN_MEMORY => 'false  
' , CACHE_BLOOMS_ON_WRITE => 'false', PREFETCH_BLOCKS_0  
N_OPEN => 'false', COMPRESSION => 'NONE', BLOCKCACHE => '  
true', BLOCKSIZE => '65536'}
```

```
1 row(s)
```

```
Took 0.1312 seconds
```

Now use SCAN to find if data exists:

```
hbase(main):004:0> scan 'iot_in'
```

As it shows, there is no data in the table:

ROW	COLUMN+CELL
-----	-------------

0 row(s)	
----------	--

Took 0.0559 seconds	
---------------------	--

```
hbase(main):005:0>
```

Populate the Table in Hive with Tez

Go back to Hive, and run below command to set Hive engine as Tez:

```
set hive.execution.engine=tez;
```

Execute below HiveQL to populate the table

```
INSERT OVERWRITE TABLE iot_data
SELECT iot_in.id, iot_in.device_id, iot_in.parameter,
iot_in.value, iot_in.datetime
FROM iot_in WHERE iot_in.device_id='SBS05';
```

Easier to copy:


```
INSERT OVERWRITE TABLE iot_data SELECT iot_in.id, iot_in.device_id, iot_in.parameter, iot_in.value, iot_in.datetime
FROM iot_in WHERE iot_in.device_id='SBS05';
```

and you should see something like this:

```
Query ID = root_20140830123030_3fee9010-e712-4c44-89ec
-1261c220e424
Total jobs = 1
Launching Job 1 out of 1
Status: Running (application id: application_140939405
7604_0003)
Map 1: -/-
Map 1: 0/1
Map 1: 0/1
Map 1: 0/1
Map 1: 0/1
Map 1: 0/1
Map 1: 0/1
Map 1: 0/1
Map 1: 1/1
Status: Finished successfully
OK
Time taken: 24.005 seconds
```

Now check the data in Hive:

```
select * from iot_data;
```

now the table should have about 85,394 rows.

Finally the HBase table (probably want to do a LIMIT):

```
hbase(main):004:0> scan 'iot_data' , {'LIMIT' => 10}
```

And get a final count on `iot_data` (in terminal):

```
[centos@ip-172-30-13-166 ~]$ hbase org.apache.hadoop.hbase  
.mapreduce.RowCounter 'iot_data'
```

Question: where is the row counter?

Improve Performance (optional)

Let's play with the data a bit to see if we can improve performance.

When populating a Hive/Hbase table the number of mappers is determined by the number of splits determined by the InputFormat used in the MapReduce job. In a typical InputFormat, it is directly proportional to the number of files and file sizes. Now suppose your HDFS block configuration is configured for 128MB(default size) and you have a files with 160MB size then it will occupy 2 block and then 2 mapper will get assigned based on the blocks. But suppose if you have 2 files with 30MB size (each file) then each file will occupy one block and

mapper will get assigned based on that.

When you are working with a large number of small files, Hive uses CombineHiveInputFormat by default. In terms of MapReduce, it ultimately translates to using CombineFileInputFormat that creates virtual splits over multiple files, grouped by common node, rack when possible. The size of the combined split is determined by

```
mapred.max.split.size  
or  
mapreduce.input.fileinputformat.split.maxsize ( in yarn/MR  
2);
```

So if you want to have less splits (less mapper action) you need to set this parameter higher.

So clear the table out:

```
truncate table iot_in;
```

Now set the performance variable:

```
0: jdbc:hive2://localhost:10000> set mapreduce.input.filei  
nputformat.split.maxsize;  
+-----+  
+  
| set
```

```
|
+-----+
+
| mapreduce.input.fileinputformat.split.maxsize=256000000
|
+-----+
+
1 row selected (0.008 seconds)
0: jdbc:hive2://localhost:10000>
```

Note the size of the max split

And then do another `LOAD DATA LOCAL ...` command.

Now another `select count(*) ...` command

What happened? Did this run go faster?

Alternative (optional)

A newer approach to loading data in Hive/HBase is to use the ImportTsv tool:

```
hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimport
tsv.separator=, -Dimporttsv.columns="HBASE_ROW_KEY,para:de
viceParameter,para:deviceValue,para:deviceId,para:dateTime
" iot_data hdfs://[master node id]:/tmp/iot_data.csv
```

Summary

Now you have seen how to hook up Hive to HBase. The tables are interchangeable. Pretty cool, huh?