

# RDBMS and Performance

MariaDB is a drop-in replacement for MySQL, the most used open source database for online applications. MariaDB falls into the category of the NewSQL products,

i.e. a product that provides unique NoSQL features together with the typical features available in relational databases.

Therefore, aspects like transaction management, durability and consistency are available together with schema or schema-less modeling,

full text storage and analysis and integration with other NoSQL technologies.

MariaDB can be part of the database infrastructure for Big Data. It is not meant to be a replacement for Hadoop, but it can be a technology used in conjunction with it. Hadoop is used in batch, ad-hoc analysis. In projects that require the processing of Terabytes or Petabytes of data,

Hadoop is definitely a good fit. The results can be queried and reported via a standard MySQL/MariaDB interface, which is compatible with virtually

all the BI tools and development frameworks available today.

Let's take a detour to check out the differences.

---

## RDBMS

So now lets take a detour and look at a RDBMS.

```
  
<h4>1. Log in to MySQL DB</h4>
```

Since it's loaded we'll use MySQL (may be MariaDB also):

```
[centos@ip-10-0-0-54 ~]$ mysql -u root -p  
Enter password:  
Welcome to the MariaDB monitor.  Commands end with ; or \g  
.  
Your MariaDB connection id is 207  
Server version: 5.5.56-MariaDB MariaDB Server  
  
Copyright (c) 2000, 2017, Oracle, MariaDB Corporation Ab and others.  
  
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.  
MariaDB [(none)]>
```



*password should be just [return]*



<h4>2. Now look at the databases</h4>

```
MariaDB [(none)]> show databases;
```

```
+-----+
```

```
| Database |
```

```
+-----+
```

```
| information_schema |
```

```
| hive |
```

```
| mysql |
```

```
| performance_schema |
```

```
| test |
```

```
+-----+
```

```
5 rows in set (0.00 sec)
```



<h4>2. And create some database, and use it</h4>

```
MariaDB [(none)]> create database class_test;
```

```
Query OK, 1 row affected (0.00 sec)
```

```
MariaDB [(none)]> use class_test;  
Database changed  
MariaDB [class_test]>
```

  
<h4>3. Let's create a table</h4>

```
MariaDB [class_test]> create table users(id int, name char  
(20), PRIMARY KEY(id));  
Query OK, 0 rows affected (0.00 sec)
```

And another table - this one will be related to the **users** table by the FOREIGN KEY **user\_id**:

```
MariaDB [class_test]> create table orders(id int, user_id  
int, order_info text, PRIMARY KEY(id), FOREIGN KEY (user_i  
d) REFERENCES users(id));  
Query OK, 0 rows affected (0.00 sec)
```



#### <h4>4. Now insert a record</h4>

```
MariaDB [class_test]> insert into users (id, name) VALUES  
(1, 'bill');  
Query OK, 1 row affected (0.00 sec)
```

Select the table:

```
MariaDB [class_test]> select * from users;  
+----+-----+  
| id | name |  
+----+-----+  
|  1 | bill |  
+----+-----+  
1 row in set (0.00 sec)
```

And we see the where clause is supported and so forth:

```
MariaDB [class_test]> select * from users where id < 2;  
+----+-----+  
| id | name |  
+----+-----+  
|  1 | bill |  
+----+-----+  
1 row in set (0.00 sec)
```

```

```

<h4>5. Create a row in the related table</h4>

```
MariaDB [class_test]> insert into orders (id, user_id, order_info) VALUES (1,1,"something");
Query OK, 1 row affected (0.00 sec)
```

And check it:

```
MariaDB [class_test]> select * from orders;
+-----+-----+-----+
| id | user_id | order_info |
+-----+-----+-----+
| 1 | 1 | something |
+-----+-----+-----+
1 row in set (0.00 sec)
```

```
MariaDB [class_test]> insert into orders (id, user_id, order_info) VALUES (2,1,"something");
Query OK, 1 row affected (0.00 sec)
```

So you see the way **2** becomes another order number, no problem:

```
MariaDB [class_test]> select * from orders;
```

```
+-----+-----+-----+
```

```
| id | user_id | order_info |
```

```
+-----+-----+-----+
```

```
| 1 | 1 | something |
```

```
| 2 | 1 | something |
```

```
+-----+-----+-----+
```

```
2 rows in set (0.00 sec)
```

But what if we do something like this:

```
MariaDB [class_test]> insert into orders (id, user_id, order_info) VALUES (1,2,"something");
```

What has happened?

## RDBMS and Transactions



<h4>6. Do the same with an active Transaction</h4>

```
MariaDB [class_test]> start transaction;
```

```
Query OK, 0 rows affected (0.00 sec)
```

Now, delete an order:

```
MariaDB [class_test]> delete from orders where id=2;  
Query OK, 1 row affected (0.01 sec)
```

And the order is now gone:

```
MariaDB [class_test]> select * from orders;  
+-----+-----+-----+  
| id | user_id | order_info |  
+-----+-----+-----+  
| 1 | 1 | something |  
+-----+-----+-----+  
1 row in set (0.00 sec)
```

But we made a mistake, so roll back the transaction:

```
MariaDB [class_test]> rollback;  
Query OK, 0 rows affected (0.00 sec)
```

```
MariaDB [class_test]> select * from orders;  
+-----+-----+-----+  
| id | user_id | order_info |  
+-----+-----+-----+  
| 1 | 1 | something |  
| 2 | 1 | something |
```




```
+-----+-----+-----+-----+
```

```
2 rows in set (0.00 sec)
```

## NOSql Performance Indicators

So now let's do some performance testing. We have a `.sql` file called `users.sql` in the directory above. :

 <https://user-images.githubusercontent.com/558905/40613898-7a6c70d6-624e-11e8-9178-7bde851ac7bd.png> align="left" width="50" height="50"

title="ToDo Logo" />

#### 7. Import into MySQL

```
[centos@ip-10-0-0-54 data]$ mysql -u root -p < users.sql
```

Now go into MySQL and see the table:

```
[centos@ip-10-0-0-54 data]$ mysql -u root -p
```

```
Enter password:
```

```
MariaDB [(none)]> show databases;
```

```
+-----+-----+-----+-----+
```

```
| Database          |
```

```
+-----+-----+-----+-----+
```

```
| information_schema |
```

```
| class_test        |
```

```
| hive |
| mysql |
| performance_schema |
| test |
| user_data |
+-----+
7 rows in set (0.00 sec)
```

```
MariaDB [(none)]> use user_data;
```

And count the table:

```
MariaDB [user_data]> select count(*) from users;
+-----+
| count(*) |
+-----+
| 1000000 |
+-----+
1 row in set (0.14 sec)
```

Rather large table, isn't it?

```

<h4>8. Look at the Data</h4>
```

```

MariaDB [user_data]> select * from users limit 1 \G
***** 1. row *****
****
      id: 1
      name: Alexandre Sporer
address: 24594 Emmitt Locks, Greenfelderview, MT 48128
      dob: 2007-05-02
      phone: 07683017318
      state: MT
1 row in set (0.00 sec)
MariaDB [user_data]>

```



*the '\G' just prints the data lengthwise*

You see that we have name, address, date of birth, phone and state in this table for 1,000,000 users.

Now see if there are any indexes on the table:

```

MariaDB [user_data]> show indexes from users \G;
***** 1. row *****
****
      Table: users
Non_unique: 0
      Key_name: PRIMARY
Seq_in_index: 1

```

Column\_name: id

Collation: A

Cardinality: 1002809

Sub\_part: NULL

Packed: NULL

Null:

Index\_type: BTREE

Comment:

Index\_comment:

1 row in set (0.00 sec)

ERROR: No query specified



*if you see more than 1 index then drop the others:*

```
MariaDB [user_data]> drop index state on users;
```

```
Query OK, 0 rows affected (0.03 sec)
```

```
Records: 0 Duplicates: 0 Warnings: 0
```

Now lets get a count by an un-indexed column:

```
MariaDB [user_data]> select count(*), state from users group by state \G
```

What were your results? We got:

...

\*\*\*\*\* 61. row \*\*\*\*\*

\*\*\*\*\*

count(\*): 15354

state: WV

\*\*\*\*\* 62. row \*\*\*\*\*

\*\*\*\*\*

count(\*): 15277

state: WY

62 rows in set (0.57 sec)

So, 0.57 seconds. Now let's run an `explain` on the query:

```
MariaDB [user_data]> explain
```

```
-> select count(*), state from users group by state \G
```

\*\*\*\*\* 1. row \*\*\*\*\*

\*\*\*\*

id: 1

select\_type: SIMPLE

table: users

type: ALL

possible\_keys: NULL

key: NULL

key\_len: NULL

ref: NULL

rows: 999129

Extra: Using temporary; Using filesort

```
1 row in set (0.00 sec)
```

So the size of the table means that **Using temporary; Using filesort** is used by MySQL to run the query.

If you add some criteria on a column in the query, this is the **EXPLAIN**:

```
MariaDB [user_data]> explain select count(*), state from u
sers where state = 'CA' group by state \G
***** 1. row *****
****

      id: 1
  select_type: SIMPLE
        table: users
         type: ALL
possible_keys: NULL
          key: NULL
       key_len: NULL
         ref: NULL
        rows: 999129
   Extra: Using where
1 row in set (0.00 sec)
```

**Using temporary; Using filesort** is gone. Now if you run the query:

```
MariaDB [user_data]> select count(*), state from users whe
re state = 'CA' group by state \G
```

```
***** 1. row *****
****
count(*): 15195
state: CA
1 row in set (0.24 sec)
```

Runtime is approximately **0.24** seconds.

```

<h4>9. Index the Column</h4>
```

Now, what if we index that column:

```
MariaDB [user_data]> alter table users add index(state);
Query OK, 0 rows affected (2.30 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

So now a b-tree structure is built to get only the rows matching the states we want.

Re-run the query:

```
MariaDB [user_data]> select count(*), state from users where
state = 'CA' group by state \G
```

```
***** 1. row *****
****
count(*): 15195
state: CA
```

What are the results now?

## Results

So we have seen where SQL databases (like MariaDB/MySQL) need to index a row by setting up an in-memory structure to make them perform better.

<https://virtuant.github.io/hadoop-overview-spark-hwx/>Go Back