

Lab: RDBMS and Performance

MariaDB is a drop-in replacement for MySQL, the most used open source database for online applications. MariaDB falls into the category of the NewSQL products,

i.e. a product that provides unique NoSQL features together with the typical features available in relational databases.

Therefore, aspects like transaction management, durability and consistency are available together with schema or schema-less modeling,

full text storage and analysis and integration with other NoSQL technologies.

MariaDB can be part of the database infrastructure for Big Data. It is not meant to be a replacement for Hadoop, but it can be a technology used in conjunction with it. Hadoop is used in batch, ad-hoc analysis. In projects that require the processing of Terabytes or Petabytes of data,

Hadoop is definitely a good fit. The results can be queried and reported via a standard MySQL/MariaDB interface, which is compatible with virtually

all the BI tools and development frameworks available today.

Let's take a detour to check out the differences.

RDBMS

So now lets take a detour and look at a RDBMS. Since it's loaded we'll use MySQL (may be MariaDB also):

```
[centos@ip-10-0-0-54 ~]$ mysql -u root -p
Enter password:
Welcome to the MariaDB monitor.  Commands end with ; or \g
.
Your MariaDB connection id is 207
Server version: 5.5.56-MariaDB MariaDB Server

Copyright (c) 2000, 2017, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.
MariaDB [(none)]>
```

Note: password should be just [return]

Now look at the databases:

```
MariaDB [(none)]> show databases;
```

```
+-----+
```

```
| Database |
```

```
+-----+
```

```
| information_schema |
```

```
| hive |
```

```
| mysql |
| performance_schema |
| test |
+-----+
5 rows in set (0.00 sec)
```

And create some database, and use it:

```
MariaDB [(none)]> create database hbase_test;
Query OK, 1 row affected (0.00 sec)

MariaDB [(none)]> use hbase_test;
Database changed
MariaDB [hbase_test]>
```

Let's create a table:

```
MariaDB [hbase_test]> create table users(id int, name char
(20), PRIMARY KEY(id));
Query OK, 0 rows affected (0.00 sec)
```

And another table - this one will be related to the `users` table by the FOREIGN KEY `user_id`:

```
MariaDB [hbase_test]> create table orders(id int, user_id
int, order_info text, PRIMARY KEY(id), FOREIGN KEY (user_i
```

```
d) REFERENCES users(id));  
Query OK, 0 rows affected (0.00 sec)
```

Now insert a record:

```
MariaDB [hbase_test]> insert into users (id, name) VALUES  
(1, 'bill');  
Query OK, 1 row affected (0.00 sec)
```

Select the table:

```
MariaDB [hbase_test]> select * from users;  
+----+-----+  
| id | name |  
+----+-----+  
|  1 | bill |  
+----+-----+  
1 row in set (0.00 sec)
```

And we see the where clause is supported and so forth:

```
MariaDB [hbase_test]> select * from users where id < 2;  
+----+-----+  
| id | name |  
+----+-----+  
|  1 | bill |
```

```
+-----+-----+
1 row in set (0.00 sec)
```

So now, create a row in the related table:

```
MariaDB [hbase_test]> insert into orders (id, user_id, order_info) VALUES (1,1,"something");
Query OK, 1 row affected (0.00 sec)
```

And check it:

```
MariaDB [hbase_test]> select * from orders;
+-----+-----+-----+
| id | user_id | order_info |
+-----+-----+-----+
| 1 | 1 | something |
+-----+-----+-----+
1 row in set (0.00 sec)
```

```
MariaDB [hbase_test]> insert into orders (id, user_id, order_info) VALUES (2,1,"something");
Query OK, 1 row affected (0.00 sec)
```

So you see the way **2** becomes another order number, no problem:

```
MariaDB [hbase_test]> select * from orders;
```

```
+-----+-----+-----+
| id | user_id | order_info |
+-----+-----+-----+
|  1 |        1 | something  |
|  2 |        1 | something  |
+-----+-----+-----+
2 rows in set (0.00 sec)
```

But what if we do something like this:

```
MariaDB [hbase_test]> insert into orders (id, user_id, order_info) VALUES (1,2,"something");
```

What has happened?

RDBMS and Transactions

This time we will do the same with an active transaction:

```
MariaDB [hbase_test]> start transaction;
Query OK, 0 rows affected (0.00 sec)
```

Now, delete an order:

```
MariaDB [hbase_test]> delete from orders where id=2;
Query OK, 1 row affected (0.01 sec)
```

And the order is now gone:

```
MariaDB [hbase_test]> select * from orders;
+-----+-----+-----+
| id | user_id | order_info |
+-----+-----+-----+
| 1 | 1 | something |
+-----+-----+-----+
1 row in set (0.00 sec)
```

But we made a mistake, so roll back the transaction:

```
MariaDB [hbase_test]> rollback;
Query OK, 0 rows affected (0.00 sec)
```

```
MariaDB [hbase_test]> select * from orders;
+-----+-----+-----+
| id | user_id | order_info |
+-----+-----+-----+
| 1 | 1 | something |
| 2 | 1 | something |
+-----+-----+-----+
2 rows in set (0.00 sec)
```

NOSql Performance Indicators

So now let's do some performance testing. We have a `.sql` file called `users.sql` in the directory above. Import into MySQL:

```
[centos@ip-10-0-0-54 data]$ mysql -u root -p < users.sql
```

Now go into MySQL and see the table:

```
[centos@ip-10-0-0-54 data]$ mysql -u root -p
```

```
Enter password:
```

```
MariaDB [(none)]> show databases;
```

```
+-----+
```

```
| Database          |
```

```
+-----+
```

```
| information_schema |
```

```
| hbase_test        |
```

```
| hive              |
```

```
| mysql             |
```

```
| performance_schema |
```

```
| test              |
```

```
| user_data         |
```

```
+-----+
```

```
7 rows in set (0.00 sec)
```

```
MariaDB [(none)]> use user_data;
```


And count the table:

```
MariaDB [user_data]> select count(*) from users;
+-----+
| count(*) |
+-----+
| 1000000  |
+-----+
1 row in set (0.14 sec)
```

Rather large table, isn't it?

Look at the data:

```
MariaDB [user_data]> select * from users limit 1 \G
***** 1. row *****
****
      id: 1
    name: Alexandre Sporer
address: 24594 Emmitt Locks, Greenfelderview, MT 48128
     dob: 2007-05-02
   phone: 07683017318
    state: MT
1 row in set (0.00 sec)
MariaDB [user_data]>
```

Note: the '\G' just prints the data lengthwise

You see that we have name, address, date of birth, phone and state in this table for 1,000,000 users.

Now see if there are any indexes on the table:

```
MariaDB [user_data]> show indexes from users \G;
***** 1. row *****
****

      Table: users
    Non_unique: 0
      Key_name: PRIMARY
Seq_in_index: 1
  Column_name: id
    Collation: A
  Cardinality: 1002809
     Sub_part: NULL
       Packed: NULL
         Null:
    Index_type: BTREE
      Comment:
    Index_comment:
1 row in set (0.00 sec)

ERROR: No query specified
```

Note: if you see more than 1 index then drop the others:

```
MariaDB [user_data]> drop index state on users;  
Query OK, 0 rows affected (0.03 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

Now lets get a count by an un-indexed column:

```
MariaDB [user_data]> select count(*), state from users gro  
up by state \G
```

What were your results? We got:

```
...  
***** 61. row *****  
*****  
count(*): 15354  
state: WV  
***** 62. row *****  
*****  
count(*): 15277  
state: WY  
62 rows in set (0.57 sec)
```

So, **0.57** seconds. Now let's run an **explain** on the query:

```

MariaDB [user_data]> explain
    -> select count(*), state from users group by state \G
***** 1. row *****
****

      id: 1
select_type: SIMPLE
      table: users
      type: ALL
possible_keys: NULL
      key: NULL
      key_len: NULL
      ref: NULL
      rows: 999129
      Extra: Using temporary; Using filesort
1 row in set (0.00 sec)

```

So the size of the table means that **Using temporary; Using filesort** is used by MySQL to run the query.

If you add some criteria on a column in the query, this is the **EXPLAIN**:

```

MariaDB [user_data]> explain select count(*), state from u
sers where state = 'CA' group by state \G
***** 1. row *****
****

      id: 1
select_type: SIMPLE

```

```
table: users
type: ALL
possible_keys: NULL
key: NULL
key_len: NULL
ref: NULL
rows: 999129
Extra: Using where
1 row in set (0.00 sec)
```

Using temporary; Using filesort is gone. Now if you run the query:

```
MariaDB [user_data]> select count(*), state from users where state = 'CA' group by state \G
***** 1. row *****
****
count(*): 15195
state: CA
1 row in set (0.24 sec)
```

Runtime is approximately 0.24 seconds.

Now, what if we index that column:

```
MariaDB [user_data]> alter table users add index(state);
Query OK, 0 rows affected (2.30 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

So now a b-tree structure is built to get only the rows matching the states we want.

Re-run the query:

```
MariaDB [user_data]> select count(*), state from users where state = 'CA' group by state \G
***** 1. row *****
****
count(*): 15195
state: CA
```

What are the results now?

Summary

So we have seen where SQL databases (like MariaDB/MySQL) need to index a row by setting up an in-memory structure to make them perform better. In the next lab we'll see how NoSQL does a similar function but this time a little differently.